

Bi/BE/CS183 2020
Prof. Matt Thomson and Lior Pachter
Problem Set 2
Due on Moodle by Thursday January 23rd at 12:00pm

Homework policy: You may collaborate with anyone on the homework and consult/use any resources. Please write down on the homework who you worked with and explain what, if any, help you received. Also list all resources that were used for the homework and explain how they helped.

Problem 1 (15 points)

Let m, x and s be fixed parameters for global alignment, and consider two sequence s_1 and s_2 that are to be aligned. Describe an algorithm that will output a random global alignment of s_1 and s_2 with probability proportional to its score.

Here m, x and s are the match, mismatch, and spacing penalties. Note there are many possible answers to this problem. A sufficient answer is one that would allow someone to completely implement the algorithm akin to following a recipe and be convinced that it should work.

To make things simpler, you can restrict all match, mismatch and spacing scores to be all negative (or all positive).

Problem 2 (10 points)

Last week we mentioned in class (and uploaded to moodle) the 2007 article *What is a gene, post-ENCODE? History and updated definition*. Read it (it's a fun read!) and answer the following:

- a) According to the article, what would be a one sentence definition of the gene for 1860s-1900s, 1910s, 1940s, 1950s, 1960s, 1970s-1980s, 1990s, 2000s?
- b) Describe how you'd think of a gene
- c) Looking at table 1 from the article, can you identify some phenomena that would be "exceptions" to your answer in b)?

Problem 3 (40 points)

In this problem you'll implement the Needleman-Wunsch algorithm. To help you with the implementation we uploaded two resources on moodle: Chapter 2 (Pairwise Alignment) of the book *Biological sequence analysis: Probabilistic models of proteins and nucleic acids* that was mentioned in class. Chapter 2 (Computation) of Lior's book *Algebraic Statistics for Computational Biology*

- a) Implement the Needleman-Wunsch algorithm, allowing for real valued match, mismatch and space parameters.
- b) Use your implementation of the Needleman-Wunsch algorithm to find all optimal *alignments* for all possible parameters for the sequences **CAGATA** and **TCAGAC**.

Clarification: When you explore the parameter space, even if you were to explore all parameter values, not all kinds of alignment would arise. For example, if you scale all scores by a constant factor, the alignment does not change. It turns out that you cannot obtain any arbitrary alignment, regardless of the scores that

you use. You can work out all possible optimal alignments by hand, but you can also use your implementation to do some brute force search over many parameters and see what alignments arise. For this question, all that you need to do is provide a list of these different alignments that you can obtain.

c) Use your algorithm to empirically determine the average length of the longest increase subsequence in a permutation of length 10.

E.g. if the permutation is 10, 2, 3, 1, 6, 7, 4, 5, 9, 8 the longest increasing subsequence has length 5 (2, 3, 4, 5, 8 or 2, 3, 6, 7, 8). You can do this by aligning the permutation with 1, 2, 3, 4, 5, 6, 7, 8, 9, 10 where a match gives score +1, a mismatch gives score 0 and a gap gives score 0, then create random permutations and measure the average score. This is not just a math exercise - the longest increasing subsequence is part of the MUMmer alignment software (<https://mummer4.github.io/>), first described in 1999. If you are curious, see the original paper *Alignment of whole genomes*, available at <https://mummer4.github.io/publications/MUMmer.pdf>

Problem 4 (15 points)

Visit the Ensembl website database page at <https://uswest.ensembl.org/info/data/ftp/index.html>. Navigate to the FTP download section and download the human cDNA (FASTA format). The file you want is named `Homo_sapiens.GRCh38.cdna.all.fa.gz` and is 65MB in size. To understand the naming convention, read the README file that is on the same FTP directory. Write some code and answer the questions below.

Hint: Biopython <https://biopython.org/> and Bioconductor <https://bioconductor.org/> are your friends. See the Bioconductor package [Biostrings](#).

- a) How many total bases of sequence are there in the transcriptome?
- b) How many transcripts are there?
- c) What is the longest transcript in the human genome?
- d) Pick your favourite random entry from the human cDNA file. Explain what each thing means.

Problem 5 (15 points)

Here we look at some samples from the 2018 article *Sirt1 protects from K-Ras-driven lung carcinogenesis*, DOI [10.15252/embr.201643879](https://doi.org/10.15252/embr.201643879). Take your time to understand what they did, and how to navigate the data repositories, because we'll revisit this and other datasets in future homeworks.

- a) Install kallisto. For instructions see <http://pachterlab.github.io/kallisto>.
- b) Download the reads for sample GSM3168965 from RNA-Seq experiment GSE115179. The GEO ("Gene Expression Omnibus") page for this experiment is <https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE115179>, but you should poke around and try to find it on your own.
- c) Quantify the experiment with kallisto and determine the most highly abundant transcript. Explain what this transcript is and what it does. Does the abundance make sense given the experiment?

Reminder: To process the sample with kallisto, first you need to use kallisto to build an index. For a given transcriptome and *kmer* size the transcriptome index needs to be built only once. The default *kmer* size is 31, and you don't need to change that. In order to build the index, first you need to download from Ensembl the appropriate transcriptome. In problem 4 you learned how to download the human transcriptome from

Ensembl. You should determine (from the paper, sample GEO page or sample SRA page) which species transcriptome to download, and then build the index. You should be able to build the index on a laptop with 8GB RAM, it usually takes 15-30 min if you have enough memory. The resulting index file will be about 2GB.

Problem 6 (5 points)

Take the Google colab notebook for HMM shown in class and extend it for 4 observed states.

Appendix

“Ah wait!”

You say, after trying to do 5(b).

“I see no download button! Not even an FTP link!”

Correct. That would be too easy.

The GEO help page (<https://www.ncbi.nlm.nih.gov/geo/info/download.html>) simply says:
All the data in GEO can be downloaded in a variety of formats using a variety of mechanisms.

In our case (and very often) the data that we want is actually hosted on the Sequence Read Archive (SRA), <https://www.ncbi.nlm.nih.gov/sra>, so the sample GSM3168965 gets another identifier associated to it. In the SRA it's known as SRR7244429. And on the sample's SRA page you can actually look at individual reads from the sequencing run (click on the **Reads** tab):
<https://trace.ncbi.nlm.nih.gov/Traces/sra/?run=SRR7244429>

But we can't download it directly. We need to use the NCBI SRA Toolkit available at:

<https://trace.ncbi.nlm.nih.gov/Traces/sra/sra.cgi?view=software>

And development information is on their github repository: <https://github.com/ncbi/sra-tools>

After the toolkit is installed, we can download the files using the tool **fastq-dump**, documented at:

<https://trace.ncbi.nlm.nih.gov/Traces/sra/sra.cgi?view=software>

Specifically, for this sample we'd use the command:

```
fastq-dump -I --split-files SRR7244429
```

Because it's a paired end read Illumina run, so it produces two **fastq** files

This downloads two files of 3.7GB each, **SRR7244429.1.fastq** and **SRR7244429.2.fastq**. You have to hope your connection doesn't get interrupted. If you are unhappy with the official documentation (as many are) here is an additional helpful resource: <https://edwards.sdsu.edu/research/fastq-dump/>

Thankfully there is a (much much) faster option: **fasterq-dump** (but it's only available for Mac and Linux)
<https://github.com/ncbi/sra-tools/wiki/HowTo:-fasterq-dump>

If you decide to use **fasterq-dump**, you can download the data with the command **fasterq-dump SRR7244429**

Finally, to help you navigate the confusing and teacherous world of getting other people's data in bioinformatics (welcome!) below is a helpful diagram of how to find the sample SRA identifier, which is what you need to run **fastq-dump** or **fasterq-dump**. Get used to this gymnastics, because you'll have to do it a few times.

[illegible]